

11/3/25

1a) A study on Cloud Analyst tools with its steps to install and run cloud analyst

1b) Write a procedure to launch virtual machine using open stack

1(a)

Cloud Analyst

It is a cloud simulator tool developed at the University of Melbourne, built on top of the CloudSim. It provides a graphical user interface to facilitate setup and visualization.

This tool helps developers and organizations optimize the deployment of applications across cloud infrastructures by analyzing factors such as request processing time, resource allocation etc..

Tools used in Cloud Analyst

1. Cloudsim Framework:

- A foundational simulation toolkit for modeling cloud computing systems
- Provides functionalities for simulating data centers, virtual machines (VMs) and cloud resources.

2. Java programming language:

- The tool is developed using Java, making it compatible with across-platform environments.
- Supports modular development & customization

3. SimJava

A discrete event simulation library used for creating models within cloud analyst.

4. GUI


Built using Java Swing, enabling users to configure simulations & visualize results easily

Advantages of cloud analyst simulator

1. Cost Effectiveness: Reduced cost and efficient resource allocation
2. Flexibility and Customization: Configurable simulations & experiment looping
3. Enhanced Decision Making: performance metrics & geographic distribution analysis
4. Educational & Research Benefits: Easy to use Interface & Rapid prototyping and testing
5. Scalability and Reliability:
Simulation of large-scale environments & dynamic simulation elements

5(b) Procedure to launch a virtual machine using openstack

1. Login To Openstack Dashboard
 - login with credentials
2. Select Project
 - select appropriate from dropdown menu
3. Check Compute Resource Quota
 - compute > overview [ensure project has sufficient resources]
to launch a VM
4. Launch instance
 - click instances & click launch instance button
5. Configure Instance Details
 - instance name, description, availability zone, count
6. Select boot source
 - choose image as boot source
7. Select ~~flavor~~ resource needs
 - CPU, RAM
8. Select Network
 - choose n/w for your VM
9. Additional settings:
 - Key Pair: select the uploaded SSH key pair for secure access
 - Click launch instance



Lab-2 25/5/23

In the given cloud simulator to host a simple web application Case

a) 6 userbases and one data centre with 50 virtual machines with each unit 1024mb memory and processor speed as 100mips.

(b) 4 userbases and 1 datacentre with 25 virtual machines each with 1024mb of memory and processor speed as 100 mips
Analyse the average minimum & maximum response time & data centre & processing time & write the findings.



2(b) Case study on workflow of amazon webservice

OUTPUT

Response time by region

| Userbase | Avg (ms) | Min (ms) | Max (ms) |
|----------|----------|----------|----------|
| VB1 | 301.173 | 225.637 | 365.204 |
| VB2 | 299.667 | 231.637 | 362.151 |
| VB3 | 299.668 | 233.141 | 366.269 |
| VB4 | 300.974 | 234.641 | 362.151 |
| VB5 | 300.984 | 225.641 | 363.651 |
| VB6 | 300.533 | 233.137 | 363.639 |


User Base Hourly Average Response Times



Data center Request Servicing Time

| Data center | Avg (ms) | Min (ms) | Max (ms) |
|-------------|----------|----------|----------|
| DC1 | 0.847 | 0.07 | 1.189 |

Data centre Loading



Cost

| Data centre | VM Cost | Data Transfer Cost | Total |
|-------------|---------|--------------------|-------|
| DC1 | 5.018 | 0.385 | 5.403 |

- 26) AWS offers a robust framework for managing workflows through its various services notably the Amazon Simple Workflow Service (SWF) and AWS Step Functions. These services help developers to build, run, and orchestrate complex applications in a scalable manner.

Simple workflow service

Purpose: coordinate distributed components of application.

Allow managing of tasks that can be executed parallelly or sequentially.

Components:

- Tasks: unit of work — activity tasks, decision tasks
- Workers - execute the tasks (like EC2)
- Domains - containers for related workflows, tasks & actors

Functionality:

- tracks the execution of workflows
- manages task assignments
- logs progress
- automatic retries for failed tasks
- good integration with other AWS services

AWS Step Functions

Purpose: higher level abstraction for building workflows
- state machines

Features

- Visual Workflow Design (GUI)
- Built In Error handling



For an order fulfillment

- Step 1: Validate Payment
Step Function triggers the lambda function to validate the payment
- Step 2: Check Inventory
Step Function triggers another lambda function to check the dynamoDB if there is enough stock
- Step 3: Shipping order
It triggers the SQS to make sure this order is put into queue & fulfilled
- Step 4: Send Notification
Triggers the amazon SNS to send the notification

1(b) Response Time By Region

| Userbase | Avg (ms) | Min (ms) | Max (ms) |
|----------|----------|----------|----------|
| UB1 | 300.264 | 232.837 | 366.338 |
| UB2 | 301.318 | 232.815 | 375.335 |
| UB3 | 300.686 | 237.311 | 360.319 |
| UB4 | 300.08 | 241.833 | 363.337 |


User Base Hourly Average Response Time



Data center Request Latency Times

| Data center | Avg (ms) | Min (ms) | Max (ms) |
|-------------|----------|----------|----------|
| DC1 | 0.549 | 0.037 | 0.856 |

~~Data center Loading~~



Cost

| Data center | VM Cost | Data Transfer Cost | Total |
|-------------|---------|--------------------|-------|
| DC1 | 2.509 | 0.255 | 2.764 |

✓
25/3

Lab program 4a) 1/4/25

In given cloud simulator to host simple web application on cloud with the configuration given below

1) Two datacenters 50 Virtual Machine each with 1024mb memory and processor speed as 100 mps, set the no. of userbase to 6. Conduct the experiment by changing different load balancing policies and conclude which load balancing policy is better by clearly stating the reason.

4b) A case study on cloud computing networking challenges & opportunities for innovation.


4a) Round Robin

Overall Response Time Summary



| | Average (ms) | Minimum (ms) | Maximum (ms) |
|------------------------------|--------------|--------------|--------------|
| Overall Response Time: | 300.91 | 227.14 | 384.64 |
| Data center Processing Time: | 1.16 | 0.07 | 1.83 |

Response Time By Region

| Userbase | Avg (ms) | Min (ms) | Max (ms) |
|----------|----------|----------|----------|
| UB1 | 300.33 | 227.141 | 379.265 |
| UB2 | 307.854 | 242.137 | 368.763 |
| UB3 | 307.884 | 239.76 | 374.761 |
| UB4 | 302.37 | 232.266 | 365.768 |
| UB5 | 300.842 | 237.636 | 370.26 |
| UB6 | 300.173 | 231.643 | 384.64 |



| Data center | Avg(ms) | min(ms) | max(ms) |
|-----------------|---------|---------|---------|
| DC ₁ | 0.852 | 0.07 | 1.186 |
| DC ₂ | 1.438 | 0.132 | 1.814 |



Cost

Total virtual machine cost: \$10.04

Total Data Transfer cost: \$0.38


Grand Total: \$10.42

| Data center | VM cost | Data Transfer Cost | Total |
|-----------------|---------|--------------------|---------|
| DC ₁ | \$5.018 | 0.188 | \$5.207 |
| DC ₂ | \$5.018 | 0.196 | \$5.215 |



Equally Spread current execution load

| | Average | Minimum | Maximum |
|-----------------------------|---------|---------|---------|
| Overall response time | 300.26 | 22.614 | 316.27 |
| Data center processing time | 1.4 | 0.07 | 1.81 |

| Userbase | Average(ms) | min(ms) | max(ms) |
|-----------------|-------------|---------|---------|
| UB ₁ | 300.785 | 225.137 | 365.151 |
| UB ₂ | 299.271 | 231.617 | 362.151 |
| UB ₃ | 299.455 | 232.641 | 366.269 |
| UB ₄ | 301.01 | 234.141 | 365.204 |
| UB ₅ | 300.423 | 235.641 | 355.643 |
| UB ₆ | 300.631 | 232.637 | 363.651 |



| Datacenter | Avg (ms) | min (ms) | Max (ms) |
|------------|----------|----------|----------|
| DC1 | 0.845 | 0.07 | 1.187 |
| DC2 | 0.449 | 0.133 | 1.814 |



Cost:

Total virtual machine cost: \$10.04

Total Data Transfer cost: \$0.38


Grand total: \$10.42

| Data center | VM cost | Data Transfer cost | Total |
|-------------|---------|--------------------|-------|
| DC1 | 0.502 | 0.193 | 0.695 |
| DC2 | 5.018 | 0.192 | 5.21 |



Throttled

| | Average | min | Max |
|-----------------------------|---------|--------|--------|
| Overall response time | 300.50 | 226.26 | 360.27 |
| Data center processing time | 1.15 | 0.07 | 1.81 |

| Uebere | avg(ms) | min(ms) | max(ms) |
|-----------------|---------|---------|---------|
| UB ₁ | 301.344 | 226.262 | 365.151 |
| UB ₂ | 299.984 | 232.262 | 361.151 |
| UB ₃ | 300.46 | 233.766 | 366.269 |
| UB ₄ | 301.809 | 235.266 | 365.204 |
| UB ₅ | 300.741 | 226.266 | 363.651 |
| UB ₆ | 300.771 | 233.132 | 364.264 |



| Datacenter | avg(ms) | min(ms) | max(ms) |
|-----------------|---------|---------|---------|
| DC ₁ | 0.849 | 0.07 | 1.187 |
| DC ₂ | 1.443 | 0.132 | 1.814 |



Cost

Total Virtual Machine cost: \$10.04

Data Transfer cost: \$0.38

Grand Total: \$10.42

| Data center | VMcost | Data transfer cost | total |
|-----------------|--------|--------------------|-------|
| DC ₁ | 5.018 | 0.193 | 5.212 |
| DC ₂ | 5.018 | 0.191 | 5.21 |

Round Robin: Distributing incoming requests evenly across all servers in a circular manner regardless of server load.

Equally spread: Distributes tasks based on current load of each server ensuring fair share of workload.

Throttled: Limits no. of requests sent to each server to prevent overloading typically by using predefined limits.

4b) Cloud computing networking challenges and opportunities for innovation

Challenges

Latency and Bandwidth: High latency and limited bandwidth impact cloud performance, especially for real-time applications.

Security risks: Data breaches and unauthorized access are major concerns in multi cloud environments.

Reliability and Downtime: Network issue can cause downtime.

Scalability: Sudden spikes can overwhelm cloud n/w.

Vendor lock in: Switching b/w cloud providers can be complex and costly due to difference in platforms, APIs & services.

Opportunities

Edge computing: Processes the data closer to the users reduces latency.

Global Accessibility: Data and services are accessible from any where with internet connectivity.

Rapid Deployment: Network services and resources can be provisioned quickly, enabling faster innovation & testing.

Enhanced cloud security features: Cloud providers offer advanced security tools (firewall, encryption, identity management) as part of their services.

Cost Efficiency: Reduces the need for physical infrastructure lowering capital expenditures and maintenance costs.

Lab program 8/4/25

3a) In a given cloud simulator to host a simulator to host a simple web application on cloud with the configuration given below

- 1) Two data centers with 50 virtual machine each with 1024mb memory and processor speed 100mbps, let the no of user bases to 6. Conduct experimentation by changing different service broker policy to analyze average minimum & maximum response.

Also analyze the total cost required to run the application

3(b) A case study on the given the web application

Ans (3a) Closest Data Center

Overall Response Time:



Data center Processing Time:

| Average(ms) | Minimum(ms) | Maximum(ms) |
|-------------|-------------|-------------|
| 300.79 | 225.64 | 366.27 |
| 1.14 | 0.07 | 1.79 |

Response Time By Region

| Userbase | Avg(ms) | Min(ms) | Max(ms) |
|----------|---------|---------|---------|
| UB1 | 301.415 | 225.637 | 365.151 |
| UB2 | 300.067 | 232.262 | 362.776 |
| UB3 | 299.902 | 233.141 | 366.269 |
| UB4 | 301.7 | 235.266 | 362.151 |
| UB5 | 300.84 | 225.641 | 364.276 |
| UB6 | 300.845 | 233.137 | 365.829 |

User Base Hourly Average Response Time




Data Center Request Servicing Times


| Data Center | Avg(ms) | Min(ms) | Max(ms) |
|-------------|---------|---------|---------|
| DC1 | 0.848 | 0.07 | 1.187 |
| DC2 | 1.442 | 0.132 | 1.792 |

Data Center Hourly Average Processing Times



Processing Time (ms)



Processing Time (ms)



Data center loading



Cost

Total Virtual Machine Cost: \$10.04
 Total Data Transfer Cost: \$0.38
 Grand Total: \$10.42

| Data Center | VM Cost | Data Transfer Cost | Total |
|-------------|---------|--------------------|-------|
| DC2 | 5.018 | 0.189 | 5.207 |
| DC1 | 5.018 | 0.196 | 5.215 |



Optimize Response Time

| | Average (ms) | Minimum (ms) | Maximum (ms) |
|-----------------------------|--------------|--------------|--------------|
| Overall Response Time: | 300.91 | 226.26 | 380.14 |
| Data Center Processing Time | 1.15 | 0.07 | 1.80 |

Response Time By Region

| Userbase | Avg (ms) | Min (ms) | Max (ms) |
|----------|----------|----------|----------|
| UB1 | 301.581 | 241.262 | 377.761 |
| UB2 | 300.366 | 243.642 | 365.764 |
| UB3 | 300.742 | 245.76 | 380.141 |
| UB4 | 300.955 | 226.263 | 363.638 |
| UB5 | 301.479 | 236.758 | 361.267 |
| UB6 | 300.318 | 238.022 | 374.143 |



User Base Hourly Average Response Times



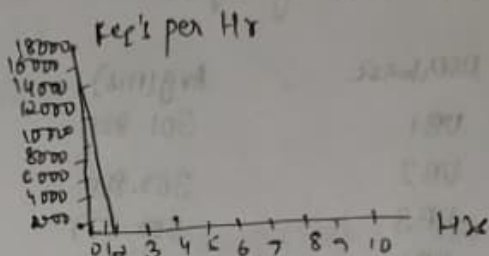

Data Center Request Servicing Times

| Data Center | Avg (ms) | Min (ms) | Max (ms) |
|-------------|----------|----------|----------|
| DC1 | 0.862 | 0.07 | 1.191 |
| DC2 | 0.1452 | 0.132 | 1.803 |

Data Center Hourly Average Processing Times



Data Center Loading



Cost

Total Virtual Machine Cost: \$10.04

Total Data Transfer Cost: \$0.38

Grand Total \$10.42

| Data Center | VM Cost | Data Transfer Cost | Total |
|-------------|---------|--------------------|-------|
| DC2 | 5.018 | 0.189 | 5.207 |
| DC1 | 5.018 | 0.196 | 5.214 |

P.T.O

Reconfigure dynamically with L



Overall Response Time Summary

| | Average (ms) | Minimum (ms) | Maximum (ms) |
|-----------------------------|--------------|--------------|--------------|
| Overall Response Time: | 301.48 | 225.71 | 366.81 |
| Data Center Processing Time | 1.82 | 0.07 | 20.25 |

Response Time By Region



| Userbase | Avg (ms) | Min (ms) | Max (ms) |
|----------|----------|----------|----------|
| UB1 | 301.801 | 225.712 | 366.814 |
| UB2 | 300.884 | 232.237 | 364.051 |
| UB3 | 300.877 | 233.641 | 366.759 |
| UB4 | 300.251 | 234.74 | 366.024 |
| UB5 | 301.354 | 226.266 | 365.12 |
| UB6 | 301.702 | 233.786 | 365.204 |

User Base Hourly Average Response Times





Data Center Request Servicing Times

| Data Center | Avg (ms) | Min (ms) | Max (ms) |
|-------------|----------|----------|----------|
| DC1 | 1.235 | 0.072 | 3.577 |
| DC2 | 2.411 | 0.135 | 20.251 |



Data Center Loading



Cost

Total Virtual Machine Cost: \$15.43

Total Data Transfer Cost: \$0.38

Grand total: \$15.82

| Data center | VM Cost | Data Transfer Cost | Total |
|-------------|---------|--------------------|-------|
| DC2 | 7.72 | 0.193 | 7.913 |
| DC1 | 7.715 | 0.192 | 7.907 |

Service Broker Policy

Description

Key Considerations

Closest Data Center

Route user workload to the data center with the least n/w latency. If multiple data centers have the same proximity it selects one randomly to balance the load.

Proximity, N/w Latency

Optimize Response Time

selects the datacenter that offers the lowest response time, considering both n/w latency & datacenter processing time

Response Time, Processing Time

Reconfigure Dynamically

Dynamically adjust the routing based on current conditions such as load & performance metrics, allowing for real time optimization

Dynamic Load Balancing, Real-time Optimization

3(b) Grep: Case Study to Find Words in Text Files

It is tool used in computers to search for specific words or phrases in text file. It helps you find information quickly without reading everything.

Scenario

You have a folder with several text files about different books, and you want to find all the books that mention the word "magic".

Steps to Use 'grep'

1. Open the Terminal
2. Go to Yo Folder (*.txt)
3. Search for "magic"
`grep "magic" *.txt`
4. See the Results

Extra Tip

- Ignore Case: MAGIC or magic

`grep -i "magic" *.txt`

- Count occurrences

`grep -c "magic" *.txt`

~~OK~~
8/9/25

5(a) Find a procedure to transfer the files from one virtual machine to another virtual machine.

(b) A case study on Performance Evaluation of Distributed File Systems in Cloud Environments.

5a) And open virtual box with specific configurations

1) General → Advanced

shared Clipboard: Bidirectional

Drag'n'Drop: Bidirectional

2) Network → Attached to: Bridged Adapter

In the terminal

- 1) sudo apt install openssh-server
- 2) sudo apt install net-tools
- 3) if config

COPY IP configuration of the guest OS

Open WINSCP

Host name: 172.25.6.249

User name: vim

password: 12345

Login ✓

Now the interface allows the transfer ^{of} the files from one virtual machine to another.

5b) Performance Metrics for Evaluation

- 1) Throughput: This measures the amount of data that can be processed in a given time frame, typically expressed in megabytes per second. High throughput is essential for applications that require large data transfers.

- 2) Latency: This is the time it takes for a request to travel from the client to the server & back. Low latency is crucial for user experience, especially in interactive applications.
- 3) Scalability: This refers to the system's ability to handle increased loads by adding more resources without significant performance degradation.
- 4) Availability: This metric indicates the system's uptime & reliability. A highly available system ensures that users can access their files whenever needed.
- 5) Consistency: This measures how up-to-date the data is across different servers. In a distributed system, ensuring that all users see the same data at the same time can be challenging.

Challenges in Performance Evaluation

Network Latency: The performance of a DFS is heavily dependent on a network speed & reliability. High latency can lead to delays in file access & retrieval.

Data Locality: When data is stored across multiple servers, accessing files that are not located close to the client can result in increased latency. Optimizing data placement is crucial for performance.

Load Balancing: Uneven distribution of data & requests across servers can lead to some servers being overloaded while others are underutilized. Effective load balancing is essential for maintaining performance.

29/4.

6(a) Docker Containers

- a) Installing Docker Engine on EC2 Instance
- b) Basic Docker commands. Understanding docker images & building images using Dockerfile. Exposing ports for web microservices, creating and using docker networks. Data persistence using docker volumes.

②

- 1) Installing Docker
- 2) Docker engine on AWS EC2 instance
- 3) Docker ⁴ images & Dockerfiles
 - a. Screenshot of C Program successfully run inside container
- 4) Exposing ports, docker networks
 - a. Sample.html showing the web page on the browser. Screenshot should show public ip of the EC2 host of docker
 - b. Screenshot of docker container running nginx (terminal of EC2 host)
 - c. Screenshot of python application successfully writing and reading from the MongoDB ~~database~~ database.
 - d. Screenshot showing mongodbs being run within network (docker command has to be clearly highlighted)
 - e. Screenshot showing python file being run within network and successfully writing and reading from MongoDB (docker command has to be clearly highlighted)
- 5) Docker Compose
 - a. Screenshot of python-mongodbs application running as a docker compose application (logs of the application)
 - b. Screenshot of 3 python application writes and reads ~~to~~ from MongoDB after scaling the python application.

6. Docker Volumes

- a. Cleaner: clearly showing command run to mount the host volume and never manually running the python application inside the container.

Docker Basic Commands

Task 1: Installing Docker Engine on EC2

1. Create EC2 instance
2. Get keypair file and SSH into the instance
3. Install docker engine on the instance
 - a. Install Docker Engine on Ubuntu
4. ~~the~~ make docker "sudo less" command
5. Verify docker engine using
docker run hello-world

Task 2: Docker images and docker files

1. Explore various images
2. The images in docker hub, need to be pulled into our EC2 instance in order to use them.
 - a. docker pull
3. To find ~~existence~~ images that exist on your instance run
 - a. docker images
4. Run the container using docker run
5. To list currently running containers use
- docker ps
6. To stop running container use docker stop
7. To remove container run docker rm

8. After you have your Dockerfile, build the image using docker build & run the container

Table 3 : Exposing ports, docker networks

1. Download sample HTML file, modify SKN
2. Create Dockerfile and then docker image having nginx: latest
 - Build docker image & tag/name it your SKN
3. Run the container. Expose the HTTP port & check connectivity

Docker

Connectivity without docker networks & how docker networks make it much easier to connect within containers.

1. Run mongod container
 - a. Use mongo: latest image
 - b. Name the container as mongod
 - c. Docker detached mode
2. Download sample.py. Modify SKN
3. Find IP using docker inspect
4. Create a Dockerfile ~~below~~

Docker Networking Options

1. Create docker bridge network called my-bridge-network
2. Run mongod container again
 - a. network: my-bridge-network
 - b. name: mongod
 - c. Exposed ports: 27017
 - d. Image: mongo: latest

Task 4: Docker Compose

1. Following files in directory
 - a. docker-compose.yml
 - b. Dockerfile
 - c. sample.py
2. Install docker compose
3. Understand syntax
4. `mn: docker-compose up`
5. scale the python application
 - a. docker-compose scale

Task 5: Docker Volumes

Volumes are the preferred mechanism for persisting data generated by and used by Docker containers.

1. Create a new directory on the EC2 host, with your ERN
 - a. sample.py
 - b. details.txt containing name, chr, section, sem
2. Add the folders inside it
3. Using Ubuntu: 18.04 as base image
4. Mount the host volume you have created & create bash shell into the container
5. Inside container, run the python program to display contents of the list